



Segmentation Models

Contents

- Unet
- Unet++
- FPN
- PSPNet
- DeepLabV3
- DeepLabV3+
- Linknet
- MAnet
- PAN
 Print to PDF ►
- UPerNet
- Segformer

[Unet](#)

[Unet++](#)

[FPN](#)

[PSPNet](#)

[DeepLabV3](#)

[DeepLabV3+](#)

[Linknet](#)

[MAnet](#)

[PAN](#)

[UPerNet](#)

[Segformer](#)

[Back to top](#)



latest

Unet

```
class segmentation_models_pytorch.Unet(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', decoder_use_batchnorm=True,
decoder_channels=(256, 128, 64, 32, 16), decoder_attention_type=None,
in_channels=3, classes=1, activation=None, aux_params=None, **kwargs) \[source\]
```

[Unet](#) is a fully convolution neural network for image semantic segmentation. Consist of *encoder* and *decoder* parts connected with *skip connections*. Encoder extract features of different spatial resolution (skip connections) which are used by decoder to define accurate segmentation mask. Use *concatenation* for fusing decoder blocks with skip connections.

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*int*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W)], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_channels** (*Tuple[int, ...]*) – List of integers which specify **in_channels** parameter for convolutions used in decoder. Length of the list should be the same as **encoder_depth**
- **decoder_use_batchnorm** (*bool*) – If **True**, BatchNorm2d layer between Conv2D and Activation layers is used. If “**inplace**” InplaceABN will be used, allows to decrease memory consumption. Available options are **True**, **False**, “**inplace**”
- **decoder_attention_type** (*str | None*) – Attention module used in decoder of the model. Available options are **None** and **scse** (<https://arxiv.org/abs/1808.08127>).
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | Callable | None*) –

 [latest](#)

An activation function to apply after the final convolution layer. Available options are “sigmoid”, “softmax”, “logsoftmax”, “tanh”, “identity”,
callable and **None**.

Default is **None**

- **aux_params** (*dict | None*) –

Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:

- classes (int): A number of classes
- pooling (str): One of “max”, “avg”. Default is “avg”
- dropout (float): Dropout factor in [0, 1)
- **activation (str): An activation function to apply “sigmoid”/“softmax”**
 (could be **None** to return logits)

- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class [`__init__\(\)`](#) function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:

`Unet`

Return type:

`torch.nn.Module`

Unet++

```
class segmentation_models_pytorch.UnetPlusPlus(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', decoder_use_batchnorm=True,
decoder_channels=(256, 128, 64, 32, 16), decoder_attention_type=None,
in_channels=3, classes=1, activation=None, aux_params=None, **kwargs) \[source\]
```

Unet++ is a fully convolution neural network for image semantic segmentation. Consist of *encoder* and *decoder* parts connected with *skip connections*. Encoder extract different spatial resolution (skip connections) which are used by decoder to define accurate segmentation mask. Decoder of Unet++ is more complex than in usual Unet.

 [latest](#)

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*int*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W),], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_channels** (*List[int]*) – List of integers which specify **in_channels** parameter for convolutions used in decoder. Length of the list should be the same as **encoder_depth**
- **decoder_use_batchnorm** (*bool*) – If **True**, BatchNorm2d layer between Conv2D and Activation layers is used. If “**inplace**” InplaceABN will be used, allows to decrease memory consumption. Available options are **True**, **False**, “**inplace**”
- **decoder_attention_type** (*str | None*) – Attention module used in decoder of the model. Available options are **None** and **scse** (<https://arxiv.org/abs/1808.08127>).
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | callable | None*) – An activation function to apply after the final convolution layer. Available options are “**sigmoid**”, “**softmax**”, “**logsoftmax**”, “**tanh**”, “**identity**”, **callable** and **None**.

Default is **None**

- **aux_params** (*dict | None*) – Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - **classes** (*int*): A number of classes
 - **pooling** (*str*): One of “max”, “avg”. Default is “avg”

 latest

- dropout (float): Dropout factor in [0, 1)
- **activation (str): An activation function to apply “sigmoid”/“softmax”**
(could be **None** to return logits)
- **kwargs (dict[str, Any])** – Arguments passed to the encoder class [`__init__\(\)`](#) function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:**Unet++****Return type:**`torch.nn.Module`**Reference:**<https://arxiv.org/abs/1807.10165>

FPN

```
class segmentation_models_pytorch.FPN(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', decoder_pyramid_channels=256,
decoder_segmentation_channels=128, decoder_merge_policy='add',
decoder_dropout=0.2, in_channels=3, classes=1, activation=None, upsampling=4,
aux_params=None, **kwargs)
```

[\[source\]](#)

[FPN](#) is a fully convolution neural network for image semantic segmentation.

Parameters:

- **encoder_name (str)** – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth (int)** – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W)], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights (str | None)** – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table  [latest](#) weights for each encoder_name)

- **decoder_pyramid_channels** (*int*) – A number of convolution filters in Feature Pyramid of [FPN](#)
- **decoder_segmentation_channels** (*int*) – A number of convolution filters in segmentation blocks of [FPN](#)
- **decoder_merge_policy** (*str*) – Determines how to merge pyramid features inside FPN. Available options are **add** and **cat**
- **decoder_dropout** (*float*) – Spatial dropout rate in range (0, 1) for feature pyramid in [FPN](#)
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | None*) –
An activation function to apply after the final convolution layer. Available options are “sigmoid”, “softmax”, “logsoftmax”, “tanh”, “identity”,
callable and **None**.

Default is **None**

- **upsampling** (*int*) – Final upsampling factor. Default is 4 to preserve input-output spatial shape identity
- **aux_params** (*dict | None*) –
Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - **classes** (*int*): A number of classes
 - **pooling** (*str*): One of “max”, “avg”. Default is “avg”
 - **dropout** (*float*): Dropout factor in [0, 1]
 - **activation (str): An activation function to apply “sigmoid”/“softmax”**
(could be **None** to return logits)

- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class  function. Applies only to [timm](#) models. Keys with [None](#) values are  latest passing.

Returns:

FPN

Return type:`torch.nn.Module`

PSPNet

```
class segmentation_models_pytorch.PSPNet(encoder_name='resnet34',
encoder_weights='imagenet', encoder_depth=3, psp_out_channels=512,
psp_use_batchnorm=True, psp_dropout=0.2, in_channels=3, classes=1,
activation=None, upsampling=8, aux_params=None, **kwargs)
```

[\[source\]](#)

PSPNet is a fully convolution neural network for image semantic segmentation. Consist of *encoder* and *Spatial Pyramid* (decoder). Spatial Pyramid build on top of encoder and does not use "fine-features" (features of high spatial resolution). PSPNet can be used for multiclass segmentation of high resolution images, however it is not good for detecting small objects and producing accurate, pixel-level mask.

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*int*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W)], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), "**imagenet**" (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **psp_out_channels** (*int*) – A number of filters in Spatial Pyramid
- **psp_use_batchnorm** (*bool*) – If **True**, BatchNorm2d layer between Conv2D and Activation layers is used. If "**inplace**" InplaceABN will be used, allows to decrease memory consumption. Available options are **True**, **False**, "**inplace**"
- **psp_dropout** (*float*) – Spatial dropout rate in [0, 1) used in Spatial Pyramid
- **in_channels** (*int*) – A number of input channels for the model, default 3 (for images)

 [latest](#)

- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | callable | None*) –
An activation function to apply after the final convolution layer. Available options are “sigmoid”, “softmax”, “logsoftmax”, “tanh”, “identity”,
callable and **None**.

Default is **None**

- **upsampling** (*int*) – Final upsampling factor. Default is 8 to preserve input-output spatial shape identity
- **aux_params** (*dict | None*) –
Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - **classes** (*int*): A number of classes
 - **pooling** (*str*): One of “max”, “avg”. Default is “avg”
 - **dropout** (*float*): Dropout factor in [0, 1)
 - **activation (str): An activation function to apply “sigmoid”/“softmax”**
(could be **None** to return logits)

- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class `__init__()` function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:

`PSPNet`

Return type:

`torch.nn.Module`

DeepLabV3

```
class segmentation_models_pytorch.DeepLabV3(encoder_name='resne'       ↗ latest
encoder_depth=5, encoder_weights='imagenet', encoder_output_stride=8,
decoder_channels=256, decoder_atrous_rates=(12, 24, 36),
```

```
decoder_aspp_separable=False, decoder_aspp_dropout=0.5, in_channels=3,
classes=1, activation=None, upsampling=None, aux_params=None, **kwargs)
```

[DeepLabV3_](#) implementation from "Rethinking Atrous Convolution for Semantic Image Segmentation" [\[source\]](#)

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*int*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W),], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), "**imagenet**" (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_channels** (*int*) – A number of convolution filters in ASPP module. Default is 256
- **encoder_output_stride** (*Literal[8, 16]*) – Downsampling factor for last encoder features (see original paper for explanation)
- **decoder_atrous_rates** (*Iterable[int]*) – Dilation rates for ASPP module (should be an iterable of 3 integer values)
- **decoder_aspp_separable** (*bool*) – Use separable convolutions in ASPP module. Default is False
- **decoder_aspp_dropout** (*float*) – Use dropout in ASPP module projection layer. Default is 0.5
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | None*) – An activation function to apply after the final convolution layer. Available options are "**sigmoid**", "**softmax**", "**logsoftmax**", "**tanh**", "**identity**",

 **callable** and **None**.

 [latest](#)

Default is **None**

- **upsampling** (*int | None*) – Final upsampling factor. Default is **None** to preserve input-output spatial shape identity
- **aux_params** (*dict | None*) –
Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - classes (int): A number of classes
 - pooling (str): One of "max", "avg". Default is "avg"
 - dropout (float): Dropout factor in [0, 1)
 - **activation (str): An activation function to apply "sigmoid"/"softmax"**
(could be **None** to return logits)
- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class [`__init__\(\)`](#) function. Applies only to [`timm`](#) models. Keys with [`None`](#) values are pruned before passing.

Returns:

DeepLabV3

Return type:

[`torch.nn.Module`](#)

[DeepLabV3+](#)

```
class segmentation_models_pytorch.DeepLabV3Plus(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', encoder_output_stride=16,
decoder_channels=256, decoder_atrous_rates=(12, 24, 36),
decoder_aspp_separable=True, decoder_aspp_dropout=0.5, in_channels=3,
classes=1, activation=None, upsampling=4, aux_params=None, **kwargs) \[source\]
```

DeepLabV3+ implementation from "Encoder-Decoder with Atrous Separable Convolution for Semantic Image Segmentation"

Parameters:

 [latest](#)

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution

- **encoder_depth** (*Literal[3, 4, 5]*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W),], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **encoder_output_stride** (*Literal[8, 16]*) – Downsampling factor for last encoder features (see original paper for explanation)
- **decoder_atrous_rates** (*Iterable[int]*) – Dilation rates for ASPP module (should be an iterable of 3 integer values)
- **decoder_aspp_separable** (*bool*) – Use separable convolutions in ASPP module. Default is True
- **decoder_aspp_dropout** (*float*) – Use dropout in ASPP module projection layer. Default is 0.5
- **decoder_channels** (*int*) – A number of convolution filters in ASPP module. Default is 256
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | None*) –

An activation function to apply after the final convolution layer. Available options are “**sigmoid**”, “**softmax**”, “**logsoftmax**”, “**tanh**”, “**identity**”,

callable and **None**.

Default is **None**

 - **upsampling** (*int*) – Final upsampling factor. Default is 4 to preserve input-output spatial shape identity.
 - **aux_params** (*dict | None*) –

Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default params:

 - **classes** (*int*): A number of classes

 latest

- pooling (str): One of "max", "avg". Default is "avg"
- dropout (float): Dropout factor in [0, 1)
- **activation (str): An activation function to apply "sigmoid"/"softmax"**
(could be **None** to return logits)
- **kwargs (dict[str, Any])** – Arguments passed to the encoder class `__init__()` function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:**DeepLabV3Plus****Return type:**`torch.nn.Module`**Reference:**<https://arxiv.org/abs/1802.02611v3>

Linknet

```
class segmentation_models_pytorch.Linknet(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', decoder_use_batchnorm=True,
in_channels=3, classes=1, activation=None, aux_params=None, **kwargs) \[source\]
```

Linknet is a fully convolution neural network for image semantic segmentation. Consist of *encoder* and *decoder* parts connected with *skip connections*. Encoder extract features of different spatial resolution (skip connections) which are used by decoder to define accurate segmentation mask. Use *sum* for fusing decoder blocks with skip connections.

 **Note**

This implementation by default has 4 skip connections (original - 3).

Parameters:

- **encoder_name (str)** – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial reso 
- **encoder_depth (int)** – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one

(e.g. for depth 0 we will have features with shapes $[(N, C, H, W),]$, for depth 1 - $[(N, C, H, W), (N, C, H // 2, W // 2)]$ and so on). Default is 5

- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_use_batchnorm** (*bool*) – If **True**, BatchNorm2d layer between Conv2D and Activation layers is used. If “**inplace**” InplaceABN will be used, allows to decrease memory consumption. Available options are **True**, **False**, “**inplace**”
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | callable | None*) –
An activation function to apply after the final convolution layer. Available options are “**sigmoid**”, “**softmax**”, “**logsoftmax**”, “**tanh**”, “**identity**”,
callable and **None**.

Default is **None**

- **aux_params** (*dict | None*) –
Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - **classes** (*int*): A number of classes
 - **pooling** (*str*): One of “max”, “avg”. Default is “avg”
 - **dropout** (*float*): Dropout factor in [0, 1)
 - **activation (str): An activation function to apply “sigmoid”/“softmax”**
(could be **None** to return logits)
- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class `__init__()` function. Applies only to `timm` models. Keys with **None** values are pruned before passing.

Returns:

Linknet

 latest

Return type:`torch.nn.Module`

MAnet

```
class segmentation_models_pytorch.MAnet(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', decoder_use_batchnorm=True,
decoder_channels=(256, 128, 64, 32, 16), decoder_pab_channels=64,
in_channels=3, classes=1, activation=None, aux_params=None, **kwargs) \[source\]
```

MAnet : Multi-scale Attention Net. The MA-Net can capture rich contextual dependencies based on the attention mechanism, using two blocks:

- Position-wise Attention Block (PAB), which captures the spatial dependencies between pixels in a global view
- Multi-scale Fusion Attention Block (MFAB), which captures the channel dependencies between any feature map by multi-scale semantic feature fusion

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*int*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W)], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_channels** (*List[int]*) – List of integers which specify **in_channels** parameter for convolutions used in decoder. Length of the list should be the same as **encoder_depth**
- **decoder_use_batchnorm** (*bool*) – If **True**, BatchNorm2d layer between Conv2D and Activation layers is used. If “**inplace**” InplaceABN will be used, allow  **latest** memory consumption. Available options are **True**, **False**, “**inplace**”

- **decoder_pab_channels** (*int*) – A number of channels for PAB module in decoder. Default is 64.
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | callable | None*) – An activation function to apply after the final convolution layer. Available options are “sigmoid”, “softmax”, “logsoftmax”, “tanh”, “identity”, **callable** and **None**.
- **aux_params** (*dict | None*) – Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - **classes** (*int*): A number of classes
 - **pooling** (*str*): One of “max”, “avg”. Default is “avg”
 - **dropout** (*float*): Dropout factor in [0, 1]
 - **activation (str): An activation function to apply “sigmoid”/“softmax”**
(could be **None** to return logits)
- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class `__init__()` function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:**MAnet****Return type:**`torch.nn.Module`

PAN

 **latest**

```
class segmentation_models_pytorch.PAN(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', encoder_output_stride=16,
```

`decoder_channels=32, in_channels=3, classes=1, activation=None, upsampling=4, aux_params=None, **kwargs)`

[\[source\]](#)

Implementation of [PAN](#) (Pyramid Attention Network).

 Note

Currently works with shape of input tensor $\geq [B \times C \times 128 \times 128]$ for pytorch $\leq 1.1.0$ and with shape of input tensor $\geq [B \times C \times 256 \times 256]$ for pytorch $= 1.3.1$

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*Literal[3, 4, 5]*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W),], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **encoder_output_stride** (*Literal[16, 32]*) – 16 or 32, if 16 use dilation in encoder last layer. Doesn’t work with ***ception***, **vgg***, **densenet*** backbones. Default is 16.
- **decoder_channels** (*int*) – A number of convolution layer filters in decoder blocks
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | Callable | None*) – An activation function to apply after the final convolution layer. Available options are “**sigmoid**”, “**softmax**”, “**logsoftmax**”, “**tanh**”, “**identity**”, **callable** and **None**.

Default is **None**

- **upsampling** (*int*) – Final upsampling factor. Default is 4 to preserve spatial shape identity

 latest

- **aux_params** (dict | None) –

Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:

- classes (int): A number of classes
- pooling (str): One of "max", "avg". Default is "avg"
- dropout (float): Dropout factor in [0, 1)
- **activation (str): An activation function to apply "sigmoid"/"softmax"**
(could be **None** to return logits)

- **kwargs** (dict[str, Any]) – Arguments passed to the encoder class `__init__()` function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:

`PAN`

Return type:

`torch.nn.Module`

UPerNet

```
class segmentation_models_pytorch.UPerNet(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet', decoder_pyramid_channels=256,
decoder_segmentation_channels=64, in_channels=3, classes=1, activation=None,
aux_params=None, **kwargs) \[source\]
```

UPerNet is a unified perceptual parsing network for image segmentation.

Parameters:

- **encoder_name** (str) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (int) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than `previous one` (e.g. for depth 0 we will have features with shapes [(N, C, H, W)], for  `C, H, W`, (N, C, H // 2, W // 2)] and so on). Default is 5

- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_pyramid_channels** (*int*) – A number of convolution filters in Feature Pyramid, default is 256
- **decoder_segmentation_channels** (*int*) – A number of convolution filters in segmentation blocks, default is 64
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | callable | None*) –
An activation function to apply after the final convolution layer. Available options are “**sigmoid**”, “**softmax**”, “**logsoftmax**”, “**tanh**”, “**identity**”,
callable and **None**.

Default is **None**

- **aux_params** (*dict | None*) –
Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:
 - **classes** (*int*): A number of classes
 - **pooling** (*str*): One of “max”, “avg”. Default is “avg”
 - **dropout** (*float*): Dropout factor in [0, 1)
 - **activation (str): An activation function to apply “sigmoid”/“softmax”**
(could be **None** to return logits)

- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class `__init__()` function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:

`UPerNet`

 [latest](#)

Return type:

`torch.nn.Module`

[Segformer](#)

```
class segmentation_models_pytorch.Segformer(encoder_name='resnet34',
encoder_depth=5, encoder_weights='imagenet',
decoder_segmentation_channels=256, in_channels=3, classes=1, activation=None,
aux_params=None, **kwargs)
```

[\[source\]](#)

Segformer is simple and efficient design for semantic segmentation with Transformers

Parameters:

- **encoder_name** (*str*) – Name of the classification model that will be used as an encoder (a.k.a backbone) to extract features of different spatial resolution
- **encoder_depth** (*int*) – A number of stages used in encoder in range [3, 5]. Each stage generate features two times smaller in spatial dimensions than previous one (e.g. for depth 0 we will have features with shapes [(N, C, H, W)], for depth 1 - [(N, C, H, W), (N, C, H // 2, W // 2)] and so on). Default is 5
- **encoder_weights** (*str | None*) – One of **None** (random initialization), “**imagenet**” (pre-training on ImageNet) and other pretrained weights (see table with available weights for each encoder_name)
- **decoder_segmentation_channels** (*int*) – A number of convolution filters in segmentation blocks, default is 256
- **in_channels** (*int*) – A number of input channels for the model, default is 3 (RGB images)
- **classes** (*int*) – A number of classes for output mask (or you can think as a number of channels of output mask)
- **activation** (*str | Callable | None*) –

An activation function to apply after the final convolution layer. Available options are “**sigmoid**”, “**softmax**”, “**logsoftmax**”, “**tanh**”, “**identity**”,

callable and **None**.

Default is **None**

- **aux_params** (*dict | None*) –

 [latest](#)

Dictionary with parameters of the auxiliary output (classification head). Auxiliary output is build on top of encoder if **aux_params** is not **None** (default). Supported params:

- classes (int): A number of classes
 - pooling (str): One of "max", "avg". Default is "avg"
 - dropout (float): Dropout factor in [0, 1)
 - **activation (str): An activation function to apply "sigmoid"/"softmax"**
(could be **None** to return logits)
- **kwargs** (*dict[str, Any]*) – Arguments passed to the encoder class [`__init__\(\)`](#) function. Applies only to `timm` models. Keys with `None` values are pruned before passing.

Returns:

Segformer

Return type:

`torch.nn.Module`

